


Deep Reinforcement Learning

Yuchu Luo

Computer Animation and Multimedia Analysis LAB

Monday, 28 August 2017






Play Ataria Game

- **Objective:** Complete the game with the highest score
- **State:** Raw pixel inputs of the game state
- **Action:** Game controls e.g. Left, Right, Up, Down
- **Reward:** Score increase/decrease at each time step

Markov Decision Process

- A set of **states** $s \in S$
- A set of **actions** (per state) $a \in A$
- A **model** $T(s,a,s')$
- A **reward function** $R(s,a,s')$
- Looking for a **policy** $\pi^*(s)$ that maximizes cumulative discounted reward: $\sum \gamma^t r_t$



Policy Learning

Find optimal policy $\pi^*(s)$

$$a \sim \pi^*(s)$$

local information

Value Learning

Find optimal Q-Value Function $Q^*(s,a)$

$$a = \arg \max_{a'} Q^*(s,a')$$

global information

$$Q^*(s_t, a_t) = \max_{\pi} E \left[\sum_{i=t}^T \gamma^{i-t} r_i \right]$$

Maximum expected future rewards starting at **state s_i** ,
choosing **action a_i** , and then following an **optimal policy π^***

Bellman Equation

$$Q^*(s, a) = E_{s' \sim \varepsilon} \left[r + \gamma \max_{a'} Q^*(s', a') \mid s, a \right]$$

Intuition:

从最佳选择的路径末端截取一小部分，余下的路径仍然是最佳路径

Solving Optimal Q-Value

Value Iteration

$$Q_{k+1}(s,a) = \mathbb{E} \left[R(s,a,s') + \gamma \max_{a'} Q_k(s',a') \mid s,a \right] = \sum_{s'} T(s,a,s') [R(s,a,s') + \gamma \max_{a'} Q_k(s',a')]$$

Convergent

Reinforcement Learning

Life is always hard — We don't know $P(s,a,s')$ and $R(s,a,s')$

Partially Observed MDP (POMDP)

Episode: sequence of states and actions

$s_0, a_0, r_0, s_1, a_1, r_1, s_2, a_2, r_2, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T, r_T$



observations (states), actions


obtain reward R

Learn to maximize the expected cumulative reward per episode

Model-Based or Model-Free?

Model-Free

Value-based



Policy-based



Recap: **Approximate Q-Learning**

Linear Value Functions

$$Q(s,a) = w_1 f_1 + w_2 f_2 + \dots + w_n f_n(s,a)$$

Feature-Based Representations

- Distance to closest ghost
- Distance to closest dot
- Number of ghosts
- $1 / (\text{dist to dot})^2$
- Is Pacman in a tunnel? (0/1)
- ...

Update

$$Q(s,a) \leftarrow (1-\alpha)Q(s,a) + \alpha(r + \gamma \max_{a'} Q(s',a'))$$

Historical experience

Learned from new (s,a,r,s') pair



$$difference = [r + \gamma \max_{a'} Q(s',a')] - Q(s,a)$$

$$Q(s,a) \leftarrow Q(s,a) + \alpha \cdot difference$$

$$w_i \leftarrow w_i + \alpha \cdot difference \cdot f_i(s,a)$$

Now, we have **deep learning**

Deep Q-Learning

$$Q(s,a;\theta) \approx Q^*(s,a)$$

Make the function approximate be a deep neural network

Loss function: $L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} [(y_i - Q(s,a;\theta_i))^2]$

Where $y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s',a';\theta_{i-1}) | s,a]$

Deep Q-Learning

Feedforward Pass



Cutest state s_t
(84x84x4) stack of last 4 frames

Convolutional
Neural
Network

Fully
Connected
Layer

Output
(4 Q-Values)

Experience Replay





Learning from batches of consecutive samples is problematic:

- Samples are correlated => inefficient learning
- current Q-network parameters determines next training samples (e.g. if maximizing action is to move left, training samples will be dominated by samples from left-hand side => can lead to bad feedback loops

Address these problems using **experience replay**

- Continually update a **replay memory** table of transitions (s_t, a_t, r_t, s_{t+1}) as game (experience) episodes are played
- Train Q-network on random mini batches of transitions from the replay memory, instead of consecutive samples
- Each transition can also contribute to multiple weight updates => greater data efficiency

Experiments



Policy Gradients

Instead of learning exact value of every (state, action) pair,
just riding the best policy from a collection of policies



left	0.6
right	0.1
fire	0.3

Probability

REINFORCE algorithm


$$\nabla_{\theta} J(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

Intuition:

- If $r(\tau)$ is high, push up the probabilities of the actions seen
- If $r(\tau)$ is low, push down the probabilities of the actions seen

Learn more in supplied materials

Actor-Critic



Neural Network

Actor Network

Critic Network

Probability

left

0.6

right

0.1

fire

0.3

left

40

right

33

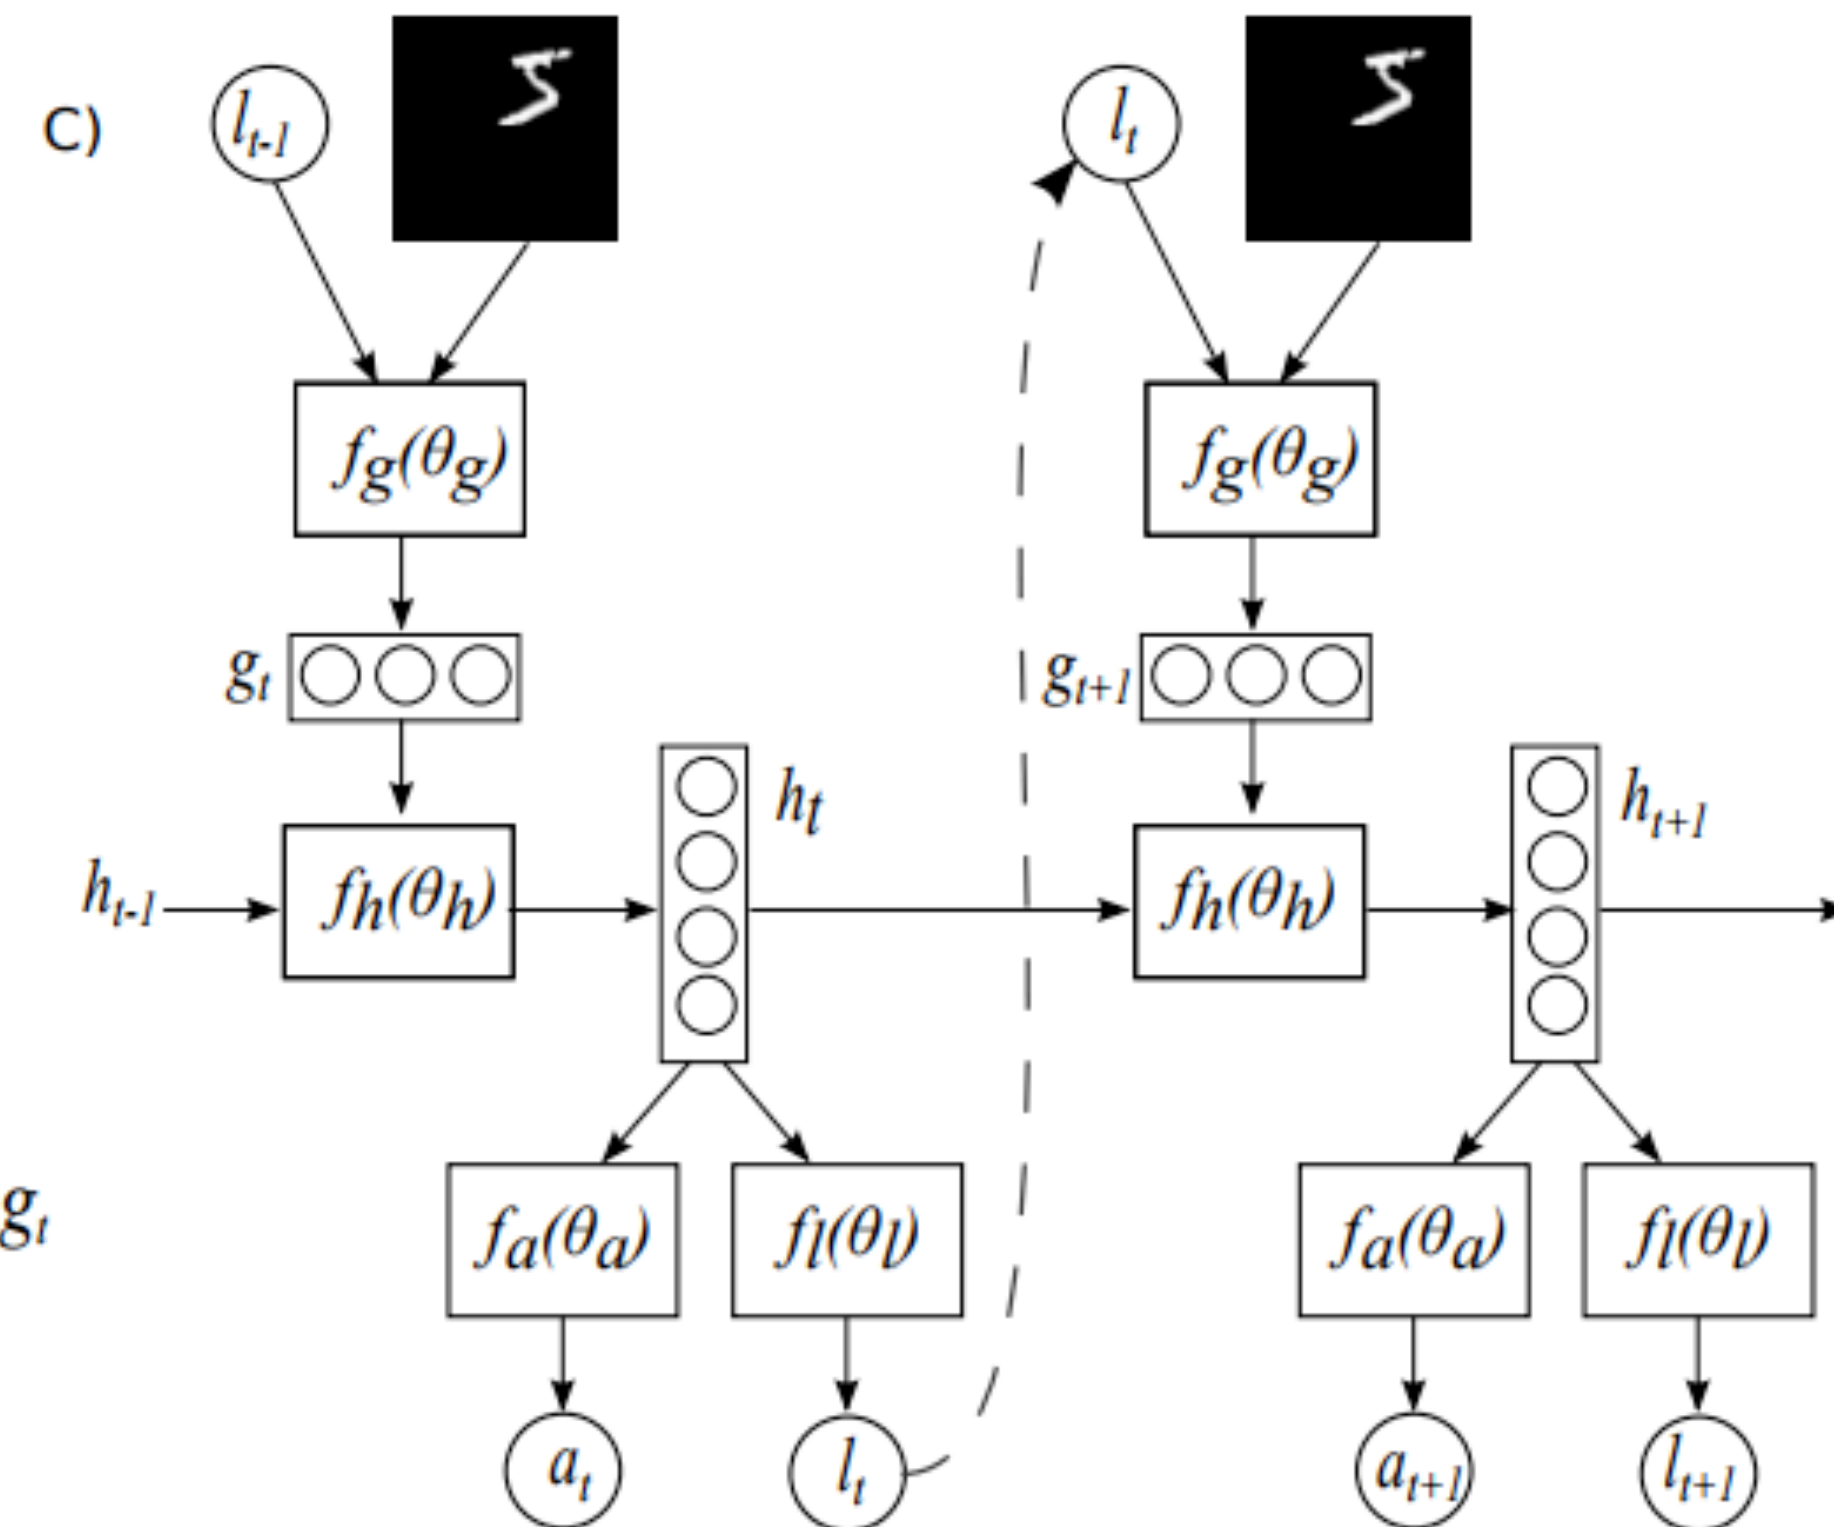
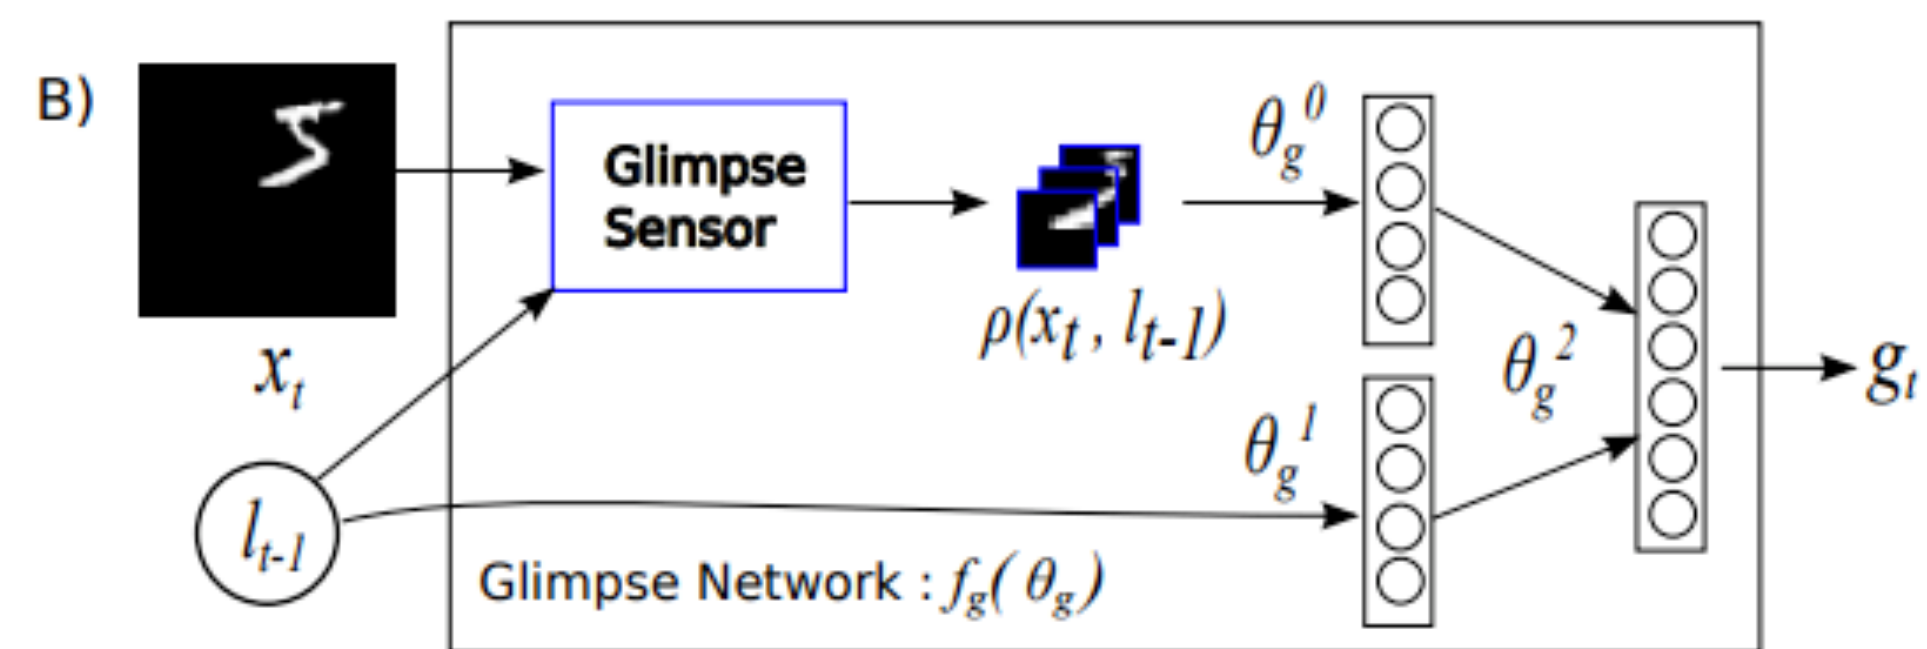

fire

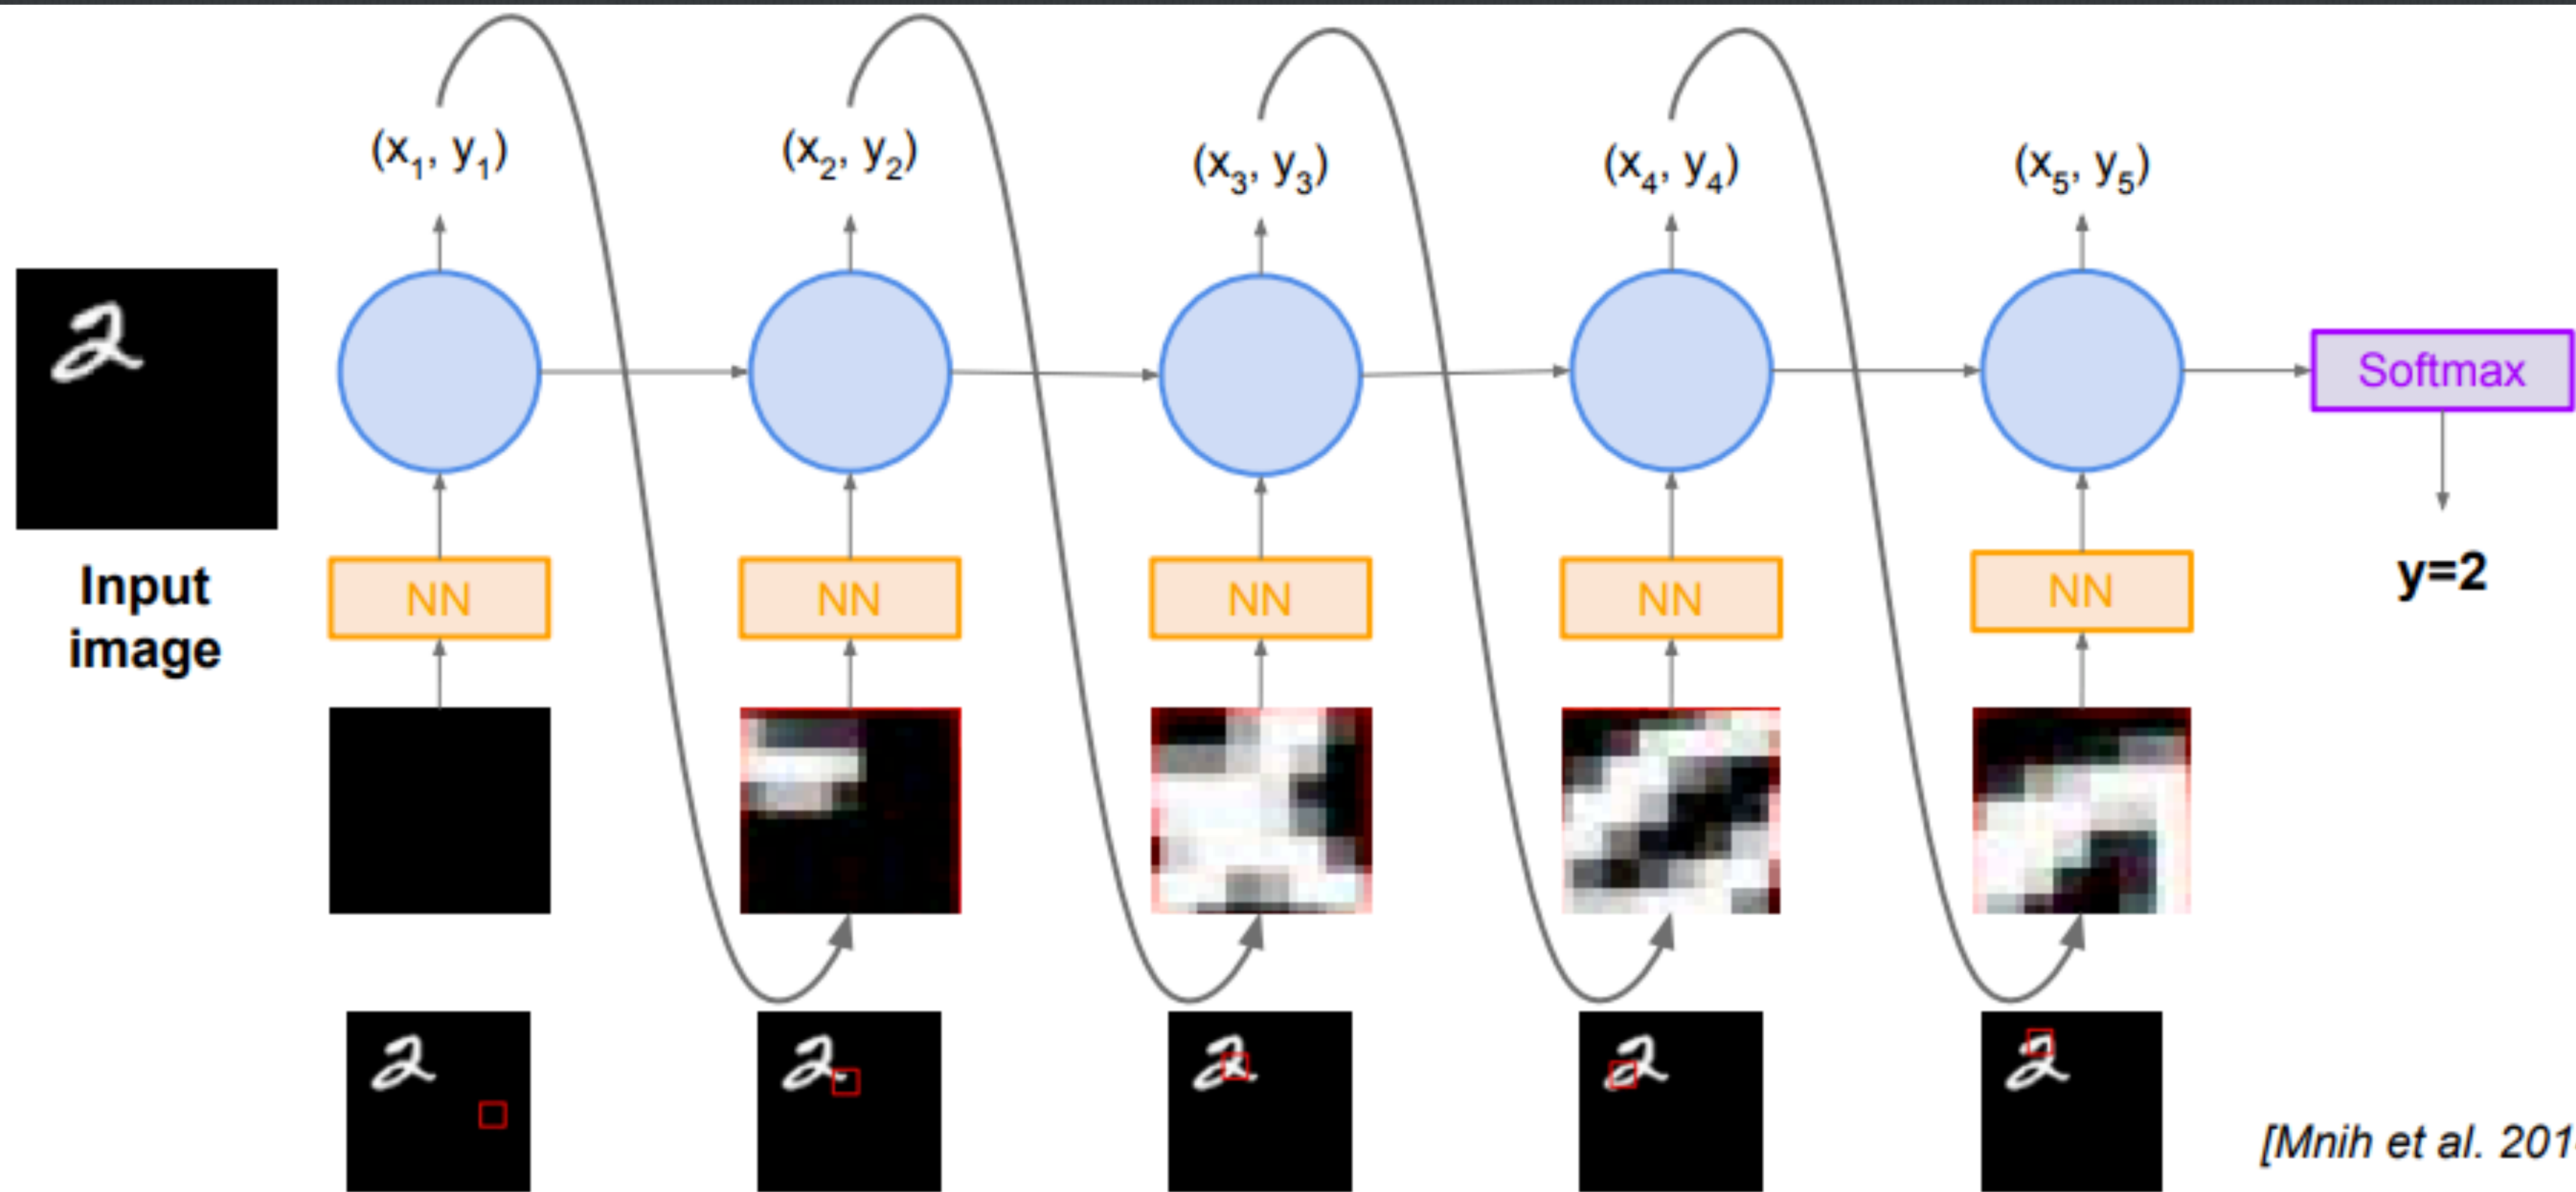
72

Q-Value

Example: Recurrent Attention Model (RAM)

Considered as a **control problem**





[Mnih et al. 2014]



Summary

- **Policy gradients**: general but suffer from high variance so requires a lot of samples. Challenge: sample-efficiency
- **Q-learning**: does not always work but when it works, usually more sample-efficient. Challenge: exploration
- Guarantees:
 - **Policy Gradients**: Converges to a local minima of $J(\theta)$, often good enough!
 - **Q-learning**: Zero guarantees since you are approximating Bellman equation with a complicated function approximator